

Name: Krish Agrawal

Roll No: 381041

Batch: A2

Assignment No: 8

Problem Statement:

Develop a text-to-speech model using Tacotron or WaveNet.

Objective:

1. To understand the fundamentals of Text-to-Speech (TTS) systems.
2. To study neural TTS architectures such as Tacotron and WaveNet.
3. To learn how text is converted into natural-sounding speech.
4. To implement a basic neural TTS pipeline using pre-trained models.
5. To evaluate speech quality and intelligibility of generated audio.

S/W Packages and H/W apparatus used:

1. Operating System: Windows/Linux/macOS
2. Kernel: Python 3.x
3. Tools: Jupyter Notebook / Anaconda
4. Hardware: CPU/GPU with minimum 4GB RAM
5. Libraries and packages used: torch, numpy, soundfile, matplotlib, TTS/torchaudio

Theory:

Text-to-Speech (TTS) systems are responsible for converting written input into audible spoken output. Conventional TTS techniques depended on concatenative or parametric approaches that often produced stilted, robotic-sounding results. In contrast, modern neural TTS solutions harness deep learning to synthesize highly expressive and natural-sounding speech directly from raw text.

Tacotron is a complete end-to-end neural TTS model that maps text input to mel-spectrograms using an encoder-decoder architecture equipped with attention mechanisms. The encoder handles the input text while the decoder produces a corresponding sequence of acoustic frames, which a vocoder subsequently converts into playable waveform audio.

WaveNet, a creation of DeepMind, is a deep generative architecture that operates directly on raw audio waveforms. By employing dilated causal convolutions, it captures long-range temporal dependencies within audio signals effectively. Although WaveNet yields extremely high-quality, lifelike speech, it demands considerable computational resources.

Neural TTS Pipeline:

Contemporary neural TTS pipelines typically involve a two-stage process: a model such as Tacotron first generates intermediate acoustic representations (mel-spectrograms), and a neural vocoder like WaveNet, WaveGlow, or HiFi-GAN then synthesizes these representations into coherent waveform audio with accurate prosody and pronunciation.

Methodology:

1. Install required libraries such as PyTorch and TTS-related packages.
2. Import necessary modules for text processing, model loading, and audio handling.
3. Load a pre-trained Tacotron or WaveNet-based TTS model.
4. Provide input text to the model.
5. Convert the text into phoneme or character embeddings.
6. Generate a mel-spectrogram representing the speech signal.
7. Use a neural vocoder to convert the spectrogram into waveform audio.
8. Save the generated speech as an audio file (e.g., .wav).
9. Play the audio to verify output quality.
10. Test the system with different input sentences.

Use Cases:

1. Virtual Assistants: Voice responses in systems like smart speakers and mobile assistants.
2. Accessibility Tools: Assists visually impaired users by reading digital text aloud.
3. Audiobook Generation: Converts written books into spoken format automatically.
4. Customer Service Automation: Interactive voice response (IVR) systems.
5. Language Learning Applications: Demonstrates pronunciation of words and sentences.

Advantages:

1. Produces natural and human-like speech.
2. Supports multiple languages and voices with appropriate models.
3. End-to-end learning reduces need for manual feature engineering.
4. Scalable for real-time applications with optimized models.
5. Can express prosody, tone, and speaking style.

Limitations:

1. Requires large datasets and computational power for training from scratch.
2. Pre-trained models may not support all languages or accents.
3. Real-time synthesis can be resource-intensive.
4. May mispronounce rare or ambiguous words.
5. Voice customization requires additional training data.

Applications:

1. Voice-enabled AI systems.
2. Navigation systems with spoken directions.
3. Assistive technology for disabilities.
4. Automated announcements and broadcasting.
5. Human-computer interaction interfaces.

Conclusion:

This assignment constructed a neural Text-to-Speech system grounded in the Tacotron or WaveNet framework. The pipeline processes input text by producing intermediate acoustic representations and subsequently converting them into waveform audio via a neural vocoder. The experiment illustrates how deep learning dramatically surpasses conventional TTS methods in producing natural, human-like speech. Systems of this kind underpin widely used voice assistants, assistive technologies, and interactive AI-driven applications.